

# Neural Network Training: A Complete Pipeline

## From Architecture Design to Performance Analysis

Machine Learning Research Group  
Computational Science Templates

November 24, 2025

### Abstract

This tutorial provides a comprehensive walkthrough of training a neural network for function approximation. We implement a multi-layer perceptron from scratch using NumPy, demonstrating forward propagation, backpropagation, and gradient descent optimization. The analysis includes architecture comparison, learning rate sensitivity, and convergence diagnostics.

## 1 Introduction

Artificial neural networks are universal function approximators capable of learning complex nonlinear mappings. This document presents a complete training pipeline, from data generation to model evaluation, with emphasis on understanding the mathematical foundations.

**Definition 1 (Feedforward Neural Network)** *A feedforward neural network is a function  $f : \mathbb{R}^n \rightarrow \mathbb{R}^m$  composed of alternating linear transformations and nonlinear activations:*

$$f(\mathbf{x}) = \sigma_L(W_L \cdot \sigma_{L-1}(W_{L-1} \cdots \sigma_1(W_1 \mathbf{x} + \mathbf{b}_1) \cdots + \mathbf{b}_{L-1}) + \mathbf{b}_L) \quad (1)$$

where  $W_l$  are weight matrices,  $\mathbf{b}_l$  are bias vectors, and  $\sigma_l$  are activation functions.

## 2 Mathematical Framework

### 2.1 Forward Propagation

For a network with  $L$  layers, the forward pass computes:

$$\mathbf{z}^{(l)} = W^{(l)} \mathbf{a}^{(l-1)} + \mathbf{b}^{(l)} \quad (\text{pre-activation}) \quad (2)$$

$$\mathbf{a}^{(l)} = \sigma(\mathbf{z}^{(l)}) \quad (\text{activation}) \quad (3)$$

## 2.2 Backpropagation

The gradient of the loss with respect to weights is computed via the chain rule:

$$\frac{\partial \mathcal{L}}{\partial W^{(l)}} = \boldsymbol{\delta}^{(l)} (\mathbf{a}^{(l-1)})^T \quad (4)$$

where the error signal propagates backward:

$$\boldsymbol{\delta}^{(l)} = (W^{(l+1)})^T \boldsymbol{\delta}^{(l+1)} \odot \sigma'(\mathbf{z}^{(l)}) \quad (5)$$

## 2.3 Activation Functions


We compare several common activation functions:

$$\text{Sigmoid: } \sigma(z) = \frac{1}{1 + e^{-z}} \quad (6)$$

$$\text{Tanh: } \sigma(z) = \frac{e^z - e^{-z}}{e^z + e^{-z}} \quad (7)$$

$$\text{ReLU: } \sigma(z) = \max(0, z) \quad (8)$$

## 3 Implementation



## 4 Training Algorithm

```

Input: Training data  $(X, y)$ , learning rate  $\eta$ , epochs  $E$ 
Output: Trained weights  $\{W^{(l)}, b^{(l)}\}$ 
Initialize weights with He initialization;
for  $epoch = 1$  to  $E$  do
    /* Forward propagation */  

     $\mathbf{a}^{(0)} \leftarrow X;$   

    for  $l = 1$  to  $L$  do  

         $\mathbf{z}^{(l)} \leftarrow W^{(l)}\mathbf{a}^{(l-1)} + \mathbf{b}^{(l)};$   

         $\mathbf{a}^{(l)} \leftarrow \sigma(\mathbf{z}^{(l)});$   

    end  

    /* Backpropagation */  

     $\delta^{(L)} \leftarrow \mathbf{a}^{(L)} - y;$   

    for  $l = L$  to  $1$  do  

         $\nabla W^{(l)} \leftarrow \delta^{(l)}(\mathbf{a}^{(l-1)})^T/m;$   

         $\nabla \mathbf{b}^{(l)} \leftarrow \text{mean}(\delta^{(l)});$   

         $\delta^{(l-1)} \leftarrow (W^{(l)})^T \delta^{(l)} \odot \sigma'(\mathbf{z}^{(l-1)});$   

    end  

    /* Gradient descent update */  

     $W^{(l)} \leftarrow W^{(l)} - \eta \nabla W^{(l)};$   

     $\mathbf{b}^{(l)} \leftarrow \mathbf{b}^{(l)} - \eta \nabla \mathbf{b}^{(l)};$   

end

```

**Algorithm 1:** Backpropagation with Gradient Descent

## 5 Results and Discussion

### 5.1 Architecture Comparison

Table 1: Neural Network Architecture Comparison

Architecture	Parameters	Final Loss	Training Time (ms)
Small	97	1.45e-01	1194.2
Medium	2241	1.38e-01	4202.3
Large	10625	1.24e-01	46278.7

The Large architecture achieved the best performance with a final MSE of 1.24e-01 using 10625 trainable parameters.

## 5.2 Observations

**Remark 1 (Capacity vs. Generalization)** *While larger networks have more representational capacity, they also require more training time and are more prone to overfitting. The gap between training and validation loss indicates generalization performance.*

**Remark 2 (Learning Rate Selection)** *The learning rate  $\eta = 0.01$  provides a good balance between convergence speed and stability. Too small ( $\eta = 0.001$ ) results in slow convergence, while too large ( $\eta = 0.1$ ) may cause oscillations or divergence.*

## 5.3 Key Findings

- Training samples: 200, Validation samples: 50
- Best architecture: Large with loss 0.1244
- The medium network [1, 64, 32, 1] offers the best trade-off between complexity and performance
- Tanh activation outperforms ReLU for this smooth target function
- He initialization is crucial for training deep networks

# 6 Limitations and Extensions

## 6.1 Current Limitations

1. **Optimization:** Plain gradient descent converges slowly. Momentum, Adam, or RMSprop would improve convergence.
2. **Regularization:** No L2 penalty or dropout is implemented, risking overfitting on larger networks.
3. **Batch Training:** Full-batch gradient descent is used; mini-batch SGD would scale better.

## 6.2 Possible Extensions

- Implement Adam optimizer with adaptive learning rates
- Add batch normalization between layers
- Implement early stopping based on validation loss
- Extend to classification with softmax output and cross-entropy loss

## 7 Conclusion

This tutorial demonstrated a complete neural network training pipeline from scratch. Key insights include the importance of architecture selection, the sensitivity to hyperparameters like learning rate, and the trade-offs between model capacity and generalization. The implementation provides a foundation for understanding more advanced deep learning frameworks.

## Further Reading

- Goodfellow, I., Bengio, Y., & Courville, A. (2016). *Deep Learning*. MIT Press.
- He, K., et al. (2015). Delving deep into rectifiers: Surpassing human-level performance on ImageNet classification.
- Kingma, D. P., & Ba, J. (2015). Adam: A method for stochastic optimization.